

Reinforcement Learning CS6700

Fall 2018

Assignment 2

Report

Author: Varun Sundar

EE16B068

24th September 2018

1 Implementation and Technical Notes

The code uses python 3.6 with sub-modules for questions. The repository adheres to the following:

- *Numpy style* documentation for the module and exposed functions.
- A *requirements.txt* for pip installing packages.
- Reproducible logs and reports.
- Code at github.com/varun19299/CS6700_Reinforcement_Learning

2 Question 1: Taxi Driver Problem

2.1 Part 1

Dynamic Programming via the compact Bellman operators was used to solve this problem.

We implement $T(J)$ by the following vectorised code:

```
np.amax(np.sum(r*P+P*np.expand_dims(J.T,2),axis=1),axis=1)
```

We get past the fact that at *town B* you cannot take action 3 (or our action 2) , by setting the rewards for action 2 (for B) as zero. Also note that since python indexing begins at zero, so do our numbering of states, stages and actions.

Where, r is the reward matrix of shape (states, states, actions); P is the probability matrix of shape (states, states, actions); J is the set of states.

We do not assume the policy to be stationary (stage independent), however, this turns out to be the case in the optimal policy.

The results may be reproduced by running:

```
python3 q1.py --stages 10
```

2.2 Part 2

The optimal policy and rewards stage wise, for $N = 10$:

```
Starting with end stage costs as [ 0.  0.  0.]
Values of J [ 16.  15.   4.5] at stage 9
Optimal policy is at stage 9 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 26.25  29.40625 18.28125] at stage 8
Optimal policy is at stage 8 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 38.265625  43.51367188 29.453125 ] at stage 7
Optimal policy is at stage 7 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 49.859375  57.30688477 41.44128418] at stage 6
Optimal policy is at stage 6 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
```

```

Values of J [ 61.65032959  70.84981537  53.24645233] at stage 5
Optimal policy is at stage 5 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 73.44839096  84.17463732  65.14957047] at stage 4
Optimal policy is at stage 4 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 85.29898071  97.31518024  77.06275252] at stage 3
Optimal policy is at stage 3 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 97.18086661 110.29839104  89.0057004 ] at stage 2
Optimal policy is at stage 2 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [109.09328351 123.1477526  100.96786822] at stage 1
Optimal policy is at stage 1 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Optimal policy is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}

```

Clearly, the optimal policy is stationary and equal to:

```

Optimal policy is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}

```

ie, ...,

- For towns A and B, Go to the nearest taxi stand and wait in line.
- For town C, Wait for a call from the dispatcher.

For $N = 20$;

```

Starting with end stage costs as [ 0.  0.  0.]
Values of J [ 16.   15.   4.5] at stage 19
Optimal policy is at stage 19 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 26.25  29.40625 18.28125] at stage 18
Optimal policy is at stage 18 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 38.265625  43.51367188 29.453125 ] at stage 17
Optimal policy is at stage 17 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 49.859375  57.30688477 41.44128418] at stage 16
Optimal policy is at stage 16 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 61.65032959  70.84981537  53.24645233] at stage 15
Optimal policy is at stage 15 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 73.44839096  84.17463732  65.14957047] at stage 14
Optimal policy is at stage 14 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 85.29898071  97.31518024  77.06275252] at stage 13

```

```

Optimal policy is at stage 13 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 97.18086661 110.29839104 89.0057004 ] at stage 12
Optimal policy is at stage 12 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 109.09328351 123.1477526 100.96786822] at stage 11
Optimal policy is at stage 11 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 121.03057586 135.88310551 112.94817246] at stage 10
Optimal policy is at stage 10 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 132.98937416 148.52138909 124.94340833] at stage 9
Optimal policy is at stage 9 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 144.96639125 161.07701436 136.9515065 ] at stage 8
Optimal policy is at stage 8 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 156.95894887 173.56225617 148.9705143 ] at stage 7
Optimal policy is at stage 7 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 168.96473159 185.9875656 160.9988241 ] at stage 6
Optimal policy is at stage 6 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 180.98177784 198.36184213 173.03505106] at stage 5
Optimal policy is at stage 5 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 193.00841445 210.69266367 185.07802059] at stage 4
Optimal policy is at stage 4 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 205.04321752 222.9864829 197.12673118] at stage 3
Optimal policy is at stage 3 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 217.08497435 235.24879433 209.18033042] at stage 2
Optimal policy is at stage 2 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 229.13265238 247.48427659 221.23809236] at stage 1
Optimal policy is at stage 1 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Optimal policy is {'state 0': 'action 1', 'state 1': 'action 1', 'state
    2': 'action 2'}

```

Again, the same policy.

Some comments on the rewards and policy:

- The rewards are unbounded - keep increasing - with stage. This is expected since there is no concept of termination in this problem.
- We donot end up with action 3 for *town B* - a sanity check.
- It is non-optimal to only go to the nearest taxi stand. We shall show this for a stage shortly.

2.3 Part 3

No, it is not optimal to go to the nearest taxi stand, irrespective of the state.

Assume this to be true. But,

$$\begin{aligned}
 J_1(C) &= \max_a E(r(C, j, a)) \\
 &= \max(2.5 + 0.5 + 0.25, 0.75 + 3 + 0.25, 3 + 1.5) \\
 &= \max(3.25, 4, 4.5)
 \end{aligned}$$

Hence, this is not so.

2.4 Code Blocks (pertinent only)

```

64 def T(J, verbose=False, stage=None):
65     '''
66     Bellman operator for maximising reward
67
68     Args:
69     * verbose: Print policy each time T is operated.
70     * stage: Which stage to operate T at.
71     '''
72     if verbose and stage:
73         policy = np.argmax(np.sum(r*P+P*J.T[np.newaxis, :, np.newaxis], axis=1), axis=1)
74         cost = np.amax(np.sum(r*P+P*J[np.newaxis, :, np.newaxis], axis=1), axis=1)
75         print(f"Policy at stage {stage} is {policy}, cost at stage {stage} is {cost}")
76     else:
77         return np.amax(np.sum(r*P+P*J[np.newaxis, :, np.newaxis], axis=1), axis=1)
78
79 def read_optimal_policy(J_opt):
80     '''
81     Prints policy for a particular optimal J.
82
83     Args:
84     * J_opt: Optimal reward.
85     '''
86     actions = np.argmax(np.sum(r*P+P*J_opt[np.newaxis, :, np.newaxis], axis=1), axis=1)
87     return {f"state {state}": f"action {action}" for state, action in zip(range(3), actions)}
```

3 Question 2

We now use similar code to question 1, except, including this into a class *Bellman* for further flexibility.

Modelling the grid-world:

- We use a 2D array, referenced by flattened indices for operating upon.
- Indices go from (0,0) to (9,9).
- Again, P , r , and J are modelled by tensors.
- *Wormholes*: Correspond to probabilities of 1 towards transition.
- *Terminal*: Collect a one time, at transit reward of +100.

Results may be replicated by:

```
88 python3 q2.py --terminal 3 --stages -1 --supress 1
```

, similar for other parameters.

3.1 Code Blocks (pertinent only)

```
89 P=np.zeros((100,100,4))
90
91 for i in range(100):
92     # Up is +10
93     if (i<90):
94         P[i,i+10,0]=0.8
95     else:
96         P[i,i,0]=0.8
97     if i%10<9:
98         P[i,i+1,0]=0.1
99     else:
100         P[i,i,0]=0.1
101     if i%10>0:
102         P[i,i-1,0]=0.1
103     else:
104         P[i,i,0]=0.1
105
106     # Down is -10 ...
107     # Left is -1 ....
108     # Right is +1 ...
109
110 # Wormholes
111 for i in [32,42,52,62]:
112     for j in range(4):
113         P[i,i+1,j]=0
114         P[i,i-1,j]=0
115         P[i,i+10,j]=0
116         P[i,i-10,j]=0
117         P[i,0,j]=1
118
119     ....
120 # Terminal stage
```

```

121
122 P[terminal,terminal,:]=1
123 if (terminal%10)<9:
124     P[terminal,terminal+1,:]=0
125 if (terminal%10)>0:
126     P[terminal,terminal-1,:]=0
127 if (terminal<90):
128     P[terminal,terminal+10,:]=0
129 if (terminal>9):
130     P[terminal,terminal-10,:]=0
131
132 r=-1*np.ones((100,100,4))
133 r[:,terminal,:]=100
134 # Collect reward only once
135 r[terminal,terminal,:]=0

```

3.2 Part 1

We stop when the maximal absolute difference ($\text{np.max}(\text{np.abs}())$) of J_i and J_{i+1} falls below a certain ϵ . For the sake of these three questions, we use $\epsilon = 1e - 6$.

The intuition for this follows from the fact that T is a contraction mapping, hence its repeated operations produce a Cauchy sequence in metric complete space. We have used to strength this fact.

3.3 Part 2 : Code Snippet

Plots on pages 10 and 11.

```

136 def plot_convergence(J_array, terminal= args.terminal, stage=0,
137     save_path=None, supress=False):
138     '''
139     Plot rewards, at stages.
140     Stages are inverted here for convinience, but nonetheless holds.
141     '''
142     J_array=np.array(J_array)
143     J_diff=np.max(np.abs(J_array[1:]-J_array[:-1]),axis=1)
144     iters=np.arange(1,len(J_diff)+1)
145     plt.plot(iters,J_diff)
146
147     plt.grid()
148
149     for j in range(len(J_diff)):
150         if (j<10 and j%3==0) or (j%10==0):
151             plt.text(j+1.15,J_diff[j]+0.15,s=f'value {J_diff[j]:.2f}')
152
153     plt.title("$max_s |J_{i+1}(s) - J_i(s)|$ vs iterations.")
154
155     # Save plot
156     if save_path:

```

```

156         plt.savefig(os.path.join(save_path,f"convergence-till-{stage}.
           png"))
157
158     if not supress:
159         plt.show()
160     else:
161         plt.close()

```

3.4 Part 3,4

We show the plots of J and *actions* as heatmaps and quiver plots respectively.

Plots on pages 12 to 17.

Comments on the plots after absolute difference convergence below a pre-determined threshold:

- Wormholes are skipped wherever they lead away from the terminal state. (1 in case of (9,9), 2 in case of (3,0)).
- Reward to go is minimal at terminal state. (since once acquired, you terminate the game).
- The general policy is to either choose an apt wormhole or the direct shortest path to the terminal state.
- Since the probability in the intended action (Up when chosen up) is dominant, we see similar actions. If this were not the case, we could see non-obvious actions.
- Colliding into the walls (and thereby retaining state) is discouraged, unless you happen to be at the terminal state.
- Thus, this policy is "greedy" and does not take into account any time-variant phenomena, memory etc. Neither does solving this given grid require this.

3.5 Code Snippet for Part 3

```

162 def quiver_actions(actions,terminal=args.terminal,stage=0, save_path=
    None,supress=False):
163     '''
164     Plot a quiver plot of the policy
165     '''
166     def _action_u(u):
167         '''
168         Horz quiver
169         -1,1 if u == left or right
170         0 else
171         '''
172         if u==2:
173             return -1

```



```



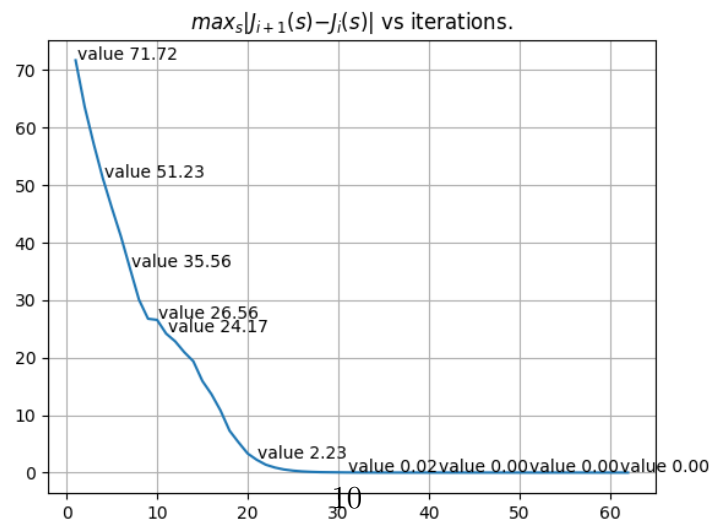
174         elif u==3:
175             return 1
176         else:
177             return 0
178     def _action_v(u):
179         '''
180         Vert quiver
181         -1,1 if u == down or up
182         0 else
183         '''
184         if u==0:
185             return 1
186         elif u==1:
187             return -1
188         else:
189             return 0
190
191     X=Y=np.arange(0.5,10.5,1)
192     U= np.array([_action_u(a) for a in actions]).reshape((10,10))
193     V= np.array([_action_v(a) for a in actions]).reshape((10,10))
194     q=plt.quiver(X,Y,U,V)
195     plt.quiverkey(q,X=8, Y=8, U=1,label='Quiver key, length = 1',
196                  labelpos='E')
197     plt.title(f"Quiver state plot at stage {stage}")
198
199     major_ticks = np.arange(0, 10, 1)
200
201     # Wormholes 1
202     for j in range(3,8):
203         plt.scatter(2.5,j+0.5,s=225,color=colors['red'])
204         plt.text(2.5, j + 0.5, 'IN1')
205     # Exit 1
206     plt.scatter(0.5,0.5,s=225,color=colors['maroon'])
207     plt.text(0.5, 0.5, 'OUT1')
208
209     # Wormholes 2
210     plt.scatter(7.5,1.5,s=225,color=colors['grey'])
211     plt.text(7.5, 1.5, 'IN2')
212     plt.scatter(7.5,9.5,s=225,color=colors['lightgrey'])
213     plt.text(7.5, 9.5, 'OUT2')
214
215     # Terminal State
216     a=args.terminal//10+0.5
217     b=args.terminal%10+0.5
218     plt.scatter(b,a,s=256,color=colors['green'])
219     plt.text(b,a, 'TERMINAL')
220
221     plt.xlim((0,10))
222     plt.ylim((0,10))

```

```
222     plt.xticks(major_ticks)
223     plt.yticks(major_ticks)
224
225     plt.grid(True)
226
227     if save_path:
228         plt.savefig(os.path.join(save_path, f"quiver-{stage}.png"))
229     if not suppress:
230         plt.show()
231     else:
232         plt.close()
```

4 References

- Classroom lectures.
- Bertsekas: *Dynamic Programming and Optimal Control, Vol 2, 3rd ed.*
- Numpy, Matplotlib Dev Documentation.

Figure 1: Convergence of J_i till $N = 10$ for **Terminal State (9,9)**Figure 2: Convergence of J_i till $N = 25$ for **Terminal State (9,9)**




Figure 4: Convergence of J_i till $N = 10$ for **Terminal State (3,0)**

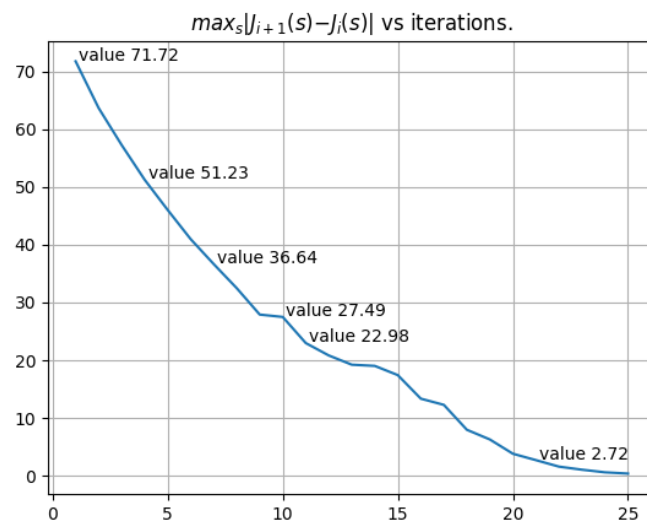
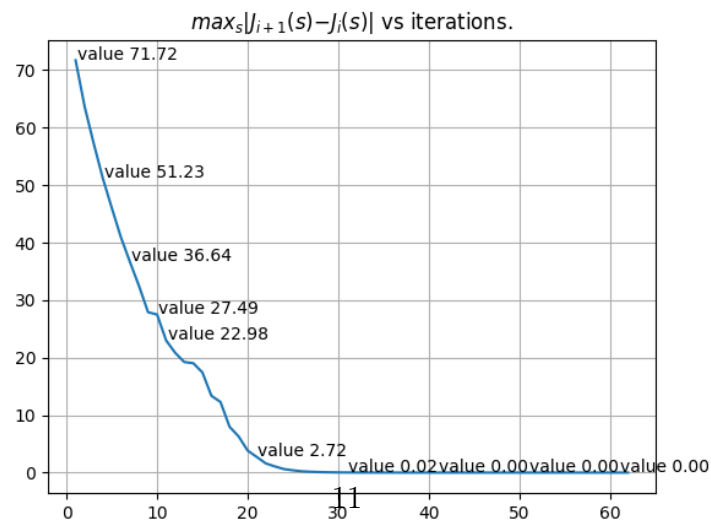








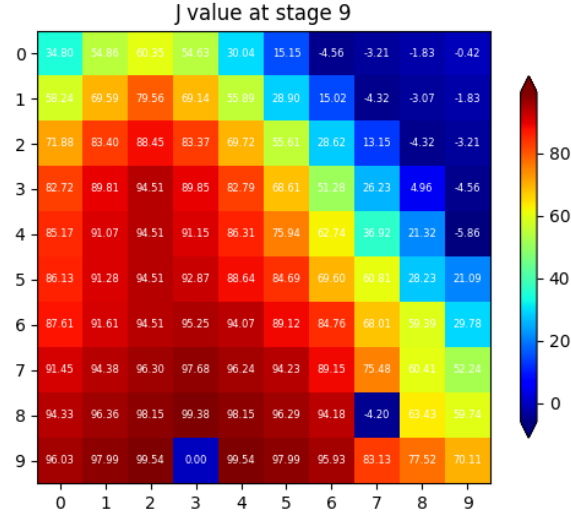

Figure 5: Convergence of J_i till $N = 25$ for **Terminal State (3,0)**


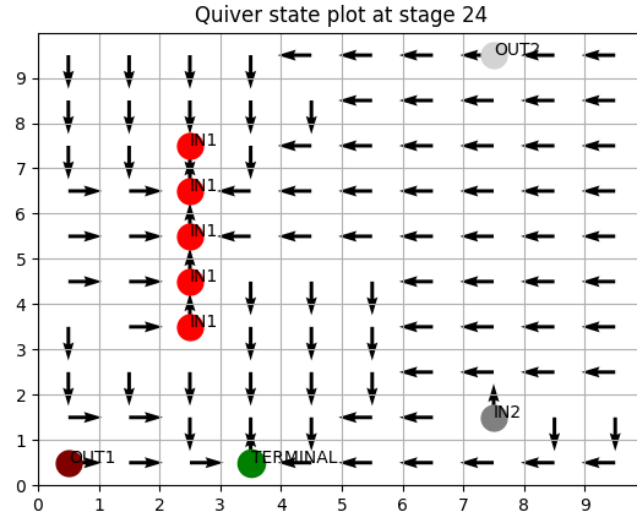




Figure 7: Heat-map of J_i till $N = 10$ for **Terminal State (9,9)**Figure 8: Quiver plot of π till $N = 10$ for **Terminal State (9,9)**

Figure 9: Heat-map of J_i till $N = 25$ for **Terminal State (9,9)**Figure 10: Quiver plot of π till $N = 25$ for **Terminal State (9,9)**

Figure 11: Heat-map of J_i till absolute difference convergence for **Terminal State (9,9)**Figure 12: Quiver plot of π till absolute difference convergence for **Terminal State (9,9)**

Figure 13: Heat-map of J_i till $N = 10$ for **Terminal State (3,0)**Figure 14: Quiver plot of π till $N = 10$ for **Terminal State (3,0)**

Figure 15: Heat-map of J_i till $N = 25$ for **Terminal State (3,0)**Figure 16: Quiver plot of π till $N = 25$ for **Terminal State (3,0)**

Figure 17: Heat-map of J_i till absolute difference convergence for **Terminal State (3,0)**Figure 18: Quiver plot of π till absolute difference convergence for **Terminal State (3,0)**