

# Reinforcement Learning CS6700

Fall 2018

## Assignment 2

### Report

Author: Varun Sundar

EE16B068

24th September 2018

# 1 Implementation and Technical Notes

The code uses python 3.6 with sub-modules for questions. The repository adheres to the following:

- *Numpy style* documentation for the module and exposed functions.
- A *requirements.txt* for pip installing packages.
- Reproducible logs and reports.

## 2 Question 1: Taxi Driver Problem

### 2.1 Part 1

Dynamic Programming via the compact Bellman operators was used to solve this problem.

We implement  $T(J)$  by the following vectorised code:

```
np.amax(np.sum(r*P+P*np.expand_dims(J.T,2),axis=1),axis=1)
```

We get past the fact that at *town B* you cannot take action 3 (or our action 2), by setting the rewards for action 2 (for B) as zero. Also note that since python indexing begins at zero, so do our numbering of states, stages and actions.

Where,  $r$  is the reward matrix of shape (states, states, actions);  $P$  is the probability matrix of shape (states, states, actions);  $J$  is the set of states.

We do not assume the policy to be stationary (stage independent), however, this turns out to be the case in the optimal policy.

The results may be reproduced by running:

```
python3 q1.py --stages 10
```

### 2.2 Part 2

The optimal policy and rewards stage wise, for  $N = 10$ :

```
Starting with end stage costs as [ 0.  0.  0.]
Values of J [ 16.  15.   4.5] at stage 9
Optimal policy is at stage 9 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 26.25  29.40625  18.28125] at stage 8
Optimal policy is at stage 8 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 38.265625  43.51367188  29.453125 ] at stage 7
Optimal policy is at stage 7 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 49.859375  57.30688477  41.44128418] at stage 6
Optimal policy is at stage 6 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 61.65032959  70.84981537  53.24645233] at stage 5
```

```

Optimal policy is at stage 5 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 73.44839096  84.17463732  65.14957047] at stage 4
Optimal policy is at stage 4 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 85.29898071  97.31518024  77.06275252] at stage 3
Optimal policy is at stage 3 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 97.18086661 110.29839104  89.0057004 ] at stage 2
Optimal policy is at stage 2 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [109.09328351 123.1477526  100.96786822] at stage 1
Optimal policy is at stage 1 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Optimal policy is {'state 0': 'action 1', 'state 1': 'action 1', 'state
    2': 'action 2'}

```

Clearly, the optimal policy is stationary and equal to:

```

Optimal policy is {'state 0': 'action 1', 'state 1': 'action 1', 'state
    2': 'action 2'}

```

ie, ...,

- For towns A and B, Go to the nearest taxi stand and wait in line.
- For town C, Wait for a call from the dispatcher.

For  $N = 20$ ;

```

Starting with end stage costs as [ 0.  0.  0.]
Values of J [ 16.   15.   4.5] at stage 19
Optimal policy is at stage 19 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 26.25   29.40625  18.28125] at stage 18
Optimal policy is at stage 18 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 38.265625  43.51367188  29.453125 ] at stage 17
Optimal policy is at stage 17 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 49.859375  57.30688477  41.44128418] at stage 16
Optimal policy is at stage 16 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 61.65032959  70.84981537  53.24645233] at stage 15
Optimal policy is at stage 15 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 73.44839096  84.17463732  65.14957047] at stage 14
Optimal policy is at stage 14 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}
Values of J [ 85.29898071  97.31518024  77.06275252] at stage 13
Optimal policy is at stage 13 is {'state 0': 'action 1', 'state 1': '
    action 1', 'state 2': 'action 2'}

```

```

Values of J [ 97.18086661 110.29839104 89.0057004 ] at stage 12
Optimal policy is at stage 12 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 109.09328351 123.1477526 100.96786822] at stage 11
Optimal policy is at stage 11 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 121.03057586 135.88310551 112.94817246] at stage 10
Optimal policy is at stage 10 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 132.98937416 148.52138909 124.94340833] at stage 9
Optimal policy is at stage 9 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 144.96639125 161.07701436 136.9515065 ] at stage 8
Optimal policy is at stage 8 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 156.95894887 173.56225617 148.9705143 ] at stage 7
Optimal policy is at stage 7 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 168.96473159 185.9875656 160.9988241 ] at stage 6
Optimal policy is at stage 6 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 180.98177784 198.36184213 173.03505106] at stage 5
Optimal policy is at stage 5 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 193.00841445 210.69266367 185.07802059] at stage 4
Optimal policy is at stage 4 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 205.04321752 222.9864829 197.12673118] at stage 3
Optimal policy is at stage 3 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 217.08497435 235.24879433 209.18033042] at stage 2
Optimal policy is at stage 2 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Values of J [ 229.13265238 247.48427659 221.23809236] at stage 1
Optimal policy is at stage 1 is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}
Optimal policy is {'state 0': 'action 1', 'state 1': 'action 1', 'state 2': 'action 2'}

```

Again, the same policy.

Some comments on the rewards and policy:

- The rewards are unbounded - keep increasing - with stage. This is expected since there is no concept of termination in this problem.
- We donot end up with action 3 for *town B* - a sanity check.
- It is non-optimal to only go to the nearest taxi stand. We shall show this for a stage shortly.

## 2.3 Part 3

No, it is not optimal to go to the nearest taxi stand, irrespective of the state.

Assume this to be true. But,

$$\begin{aligned}
 J_1(C) &= \max_a E(r(C, j, a)) \\
 &= \max(2.5 + 0.5 + 0.25, 0.75 + 3 + 0.25, 3 + 1.5) \\
 &= \max(3.25, 4, 4.5)
 \end{aligned}$$

Hence, this is not so.

## 2.4 Code Blocks (pertinent only)

```

64 def T(J, verbose=False, stage=None):
65     '''
66     Bellman operator for maximising reward
67
68     Args:
69     * verbose: Print policy each time T is operated.
70     * stage: Which stage to operate T at.
71     '''
72     if verbose and stage:
73         policy = np.argmax(np.sum(r*P+P*J.T[np.newaxis, :, np.newaxis], axis=1), axis=1)
74         cost = np.amax(np.sum(r*P+P*J[np.newaxis, :, np.newaxis], axis=1), axis=1)
75         print(f"Policy at stage {stage} is {policy}, cost at stage {cost}")
76     else:
77         return np.amax(np.sum(r*P+P*J[np.newaxis, :, np.newaxis], axis=1), axis=1)
78
79 def read_optimal_policy(J_opt):
80     '''
81     Prints policy for a particular optimal J.
82
83     Args:
84     * J_opt: Optimal reward.
85     '''
86     actions = np.argmax(np.sum(r*P+P*J_opt[np.newaxis, :, np.newaxis], axis=1), axis=1)
87     return {f"state {state}": f"action {action}" for state, action in zip(range(3), actions)}

```

## 3 Question 2

We now use similar code to question 1, except, including this into a class *Bellman* for further flexibility.

Modelling the grid-world:

- We use a 2D array, referenced by flattened indices for operating upon.
- Indices go from (0,0) to (9,9).
- Again,  $P$ ,  $r$ , and  $J$  are modelled by tensors.
- *Wormholes*: Correspond to probabilities of 1 towards transition.
- *Terminal*: Collect a one time, at transit reward of +100.

### 3.1 Code Blocks (pertinent only)

```

88 P=np.zeros((100,100,4))
89
90 for i in range(100):
91     # Up is +10
92     if (i<90):
93         P[i,i+10,0]=0.8
94     else:
95         P[i,i,0]=0.8
96     if i%10<9:
97         P[i,i+1,0]=0.1
98     else:
99         P[i,i,0]=0.1
100     if i%10>0:
101         P[i,i-1,0]=0.1
102     else:
103         P[i,i,0]=0.1
104
105     # Down is -10 ...
106     # Left is -1 ....
107     # Right is +1 ...
108
109 # Wormholes
110 for i in [32,42,52,62]:
111     for j in range(4):
112         P[i,i+1,j]=0
113         P[i,i-1,j]=0
114         P[i,i+10,j]=0
115         P[i,i-10,j]=0
116         P[i,0,j]=1
117
118     ....
119 # Terminal stage
120
121 P[terminal,terminal,:]=1
122 if (terminal%10)<9:
123     P[terminal,terminal+1,:]=0
124 if (terminal%10)>0:

```

```

125     P[terminal,terminal-1,:]=0
126 if (terminal<90):
127     P[terminal,terminal+10,:]=0
128 if (terminal>9):
129     P[terminal,terminal-10,:]=0
130
131 r=-1*np.ones((100,100,4))
132 r[:,terminal,:]=100
133 # Collect reward only once
134 r[terminal,terminal,:]=0

```

## 3.2 Part 1

We stop when the maximal absolute difference ( $\max(\text{np.abs}())$ ) of  $J_i$  and  $J_{i+1}$  falls below a certain  $\epsilon$ . For the sake of these three questions, we use  $\epsilon = 1e - 6$ .

The intuition for this follows from the fact that  $T$  is a contraction mapping, hence its repeated operations produce a Cauchy sequence in metric complete space. We have used to strength this fact.

## 3.3 Part 2 : Code Snippet

Plots on pages 10 and 11.

```

135 def plot_convergence(J_array, terminal= args.terminal, stage=0,
136     save_path=None, supress=False):
137     '''
138     Plot rewards, at stages.
139     Stages are inverted here for convinience, but nonetheless holds.
140     '''
141     J_array=np.array(J_array)
142     J_diff=np.max(np.abs(J_array[1:]-J_array[:-1]),axis=1)
143     iters=np.arange(1,len(J_diff)+1)
144     plt.plot(iters,J_diff)
145
146     plt.grid()
147
148     for j in range(len(J_diff)):
149         if (j<10 and j%3==0) or (j%10==0):
150             plt.text(j+1.15,J_diff[j]+0.15,s=f'value {J_diff[j]:.2f}')
151
152     plt.title("$max_s |J_{i+1}(s) - J_i(s)|$ vs iterations.")
153
154     # Save plot
155     if save_path:
156         plt.savefig(os.path.join(save_path,f"convergence-till-{stage}.png"))
157
158     if not supress:
159         plt.show()

```

```

159     else:
160         plt.close()

```

### 3.4 Part 3,4

We show the plots of  $J$  and *actions* as heatmaps and quiver plots respectively.

Plots on pages 12 to 17.

Comments on the plots after absolute difference convergence below a pre-determined threshold:

- Wormholes are skipped wherever they lead away from the terminal state. (1 in case of (9,9), 2 in case of (3,0)).
- Reward to go is minimal at terminal state. (since once acquired, you terminate the game).
- The general policy is to either choose an apt wormhole or the direct shortest path to the terminal state.
- Since the probability in the intended action (Up when chosen up) is dominant, we see similar actions. If this were not the case, we could see non-obvious actions.
- Colliding into the walls (and thereby retaining state) is discouraged, unless you happen to be at the terminal state.
- Thus, this policy is "greedy" and does not take into account any time-variant phenomena, memory etc. Neither does solving this given grid require this.

### 3.5 Code Snippet for Part 3

```

161 def quiver_actions(actions,terminal=args.terminal,stage=0, save_path=
    None,supress=False):
162     '''
163     Plot a quiver plot of the policy
164     '''
165     def _action_u(u):
166         '''
167         Horz quiver
168         -1,1 if u == left or right
169         0 else
170         '''
171         if u==2:
172             return -1
173         elif u==3:
174             return 1
175         else:
176             return 0
177     def _action_v(u):

```



```



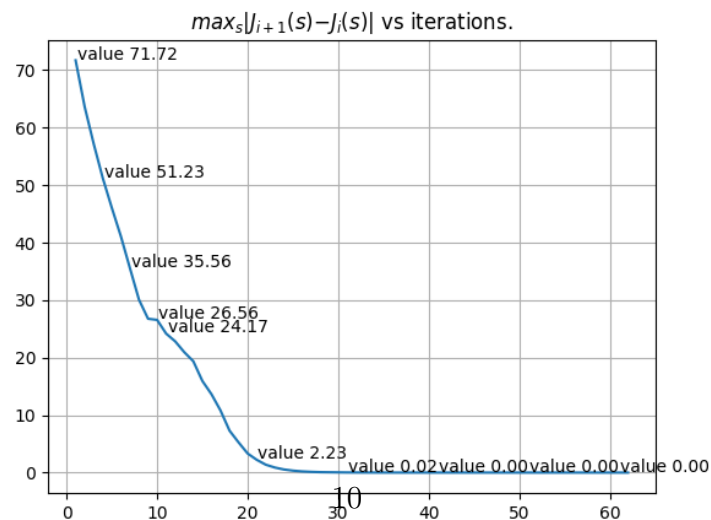
178     '''
179     Vert quiver
180     -1,1 if u == down or up
181     0 else
182     '''
183     if u==0:
184         return 1
185     elif u==1:
186         return -1
187     else:
188         return 0
189
190     X=Y=np.arange(0.5,10.5,1)
191     U= np.array([_action_u(a) for a in actions]).reshape((10,10))
192     V= np.array([_action_v(a) for a in actions]).reshape((10,10))
193     q=plt.quiver(X,Y,U,V)
194     plt.quiverkey(q,X=8, Y=8, U=1,label='Quiver key, length = 1',
195                  labelpos='E')
196     plt.title(f"Quiver state plot at stage {stage}")
197
198     major_ticks = np.arange(0, 10, 1)
199
200     # Wormholes 1
201     for j in range(3,8):
202         plt.scatter(2.5,j+0.5,s=225,color=colors['red'])
203         plt.text(2.5, j + 0.5, 'IN1')
204     # Exit 1
205     plt.scatter(0.5,0.5,s=225,color=colors['maroon'])
206     plt.text(0.5, 0.5, 'OUT1')
207
208     # Wormholes 2
209     plt.scatter(7.5,1.5,s=225,color=colors['grey'])
210     plt.text(7.5, 1.5, 'IN2')
211     plt.scatter(7.5,9.5,s=225,color=colors['lightgrey'])
212     plt.text(7.5, 9.5, 'OUT2')
213
214     # Terminal State
215     a=args.terminal//10+0.5
216     b=args.terminal%10+0.5
217     plt.scatter(b,a,s=256,color=colors['green'])
218     plt.text(b,a, 'TERMINAL')
219
220     plt.xlim((0,10))
221     plt.ylim((0,10))
222     plt.xticks(major_ticks)
223     plt.yticks(major_ticks)
224
225     plt.grid(True)


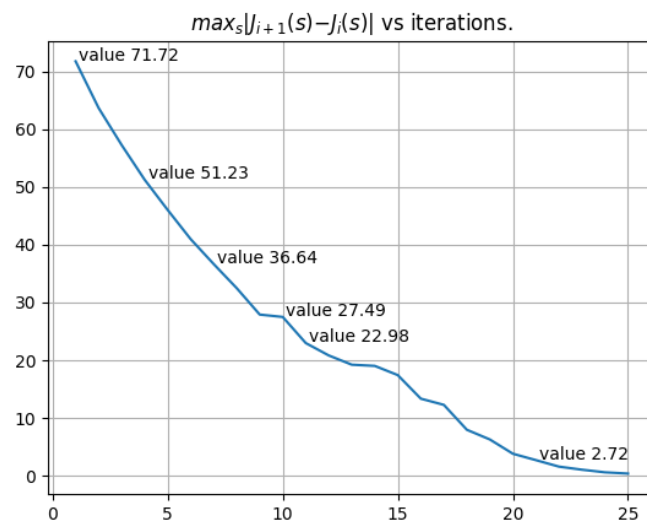
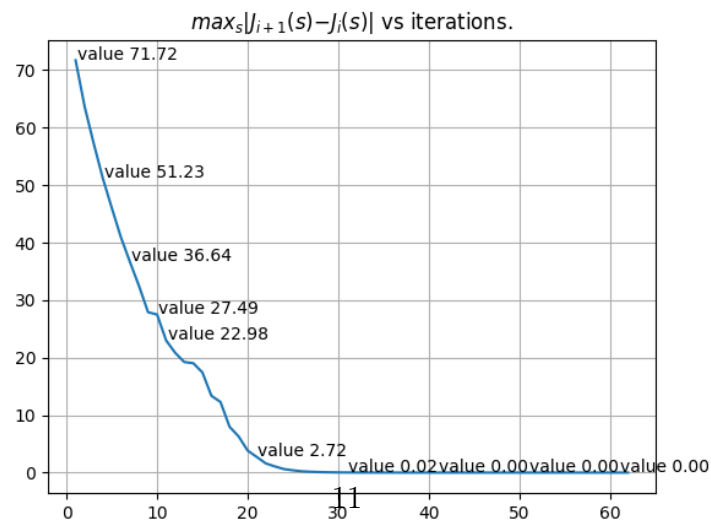
```

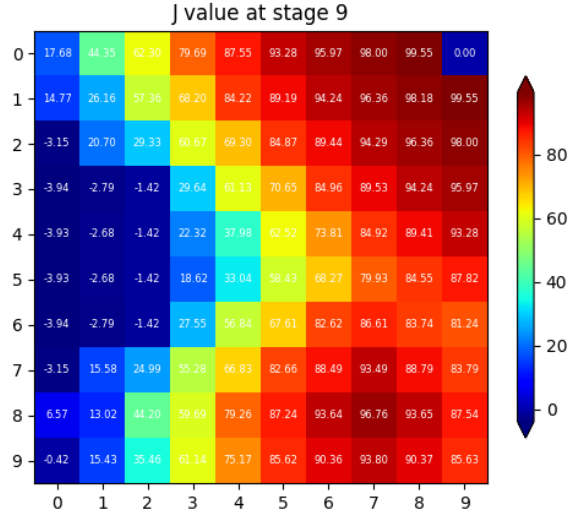

```
226     if save_path:
227         plt.savefig(os.path.join(save_path, f"quiver-{stage}.png"))
228     if not suppress:
229         plt.show()
230     else:
231         plt.close()
```


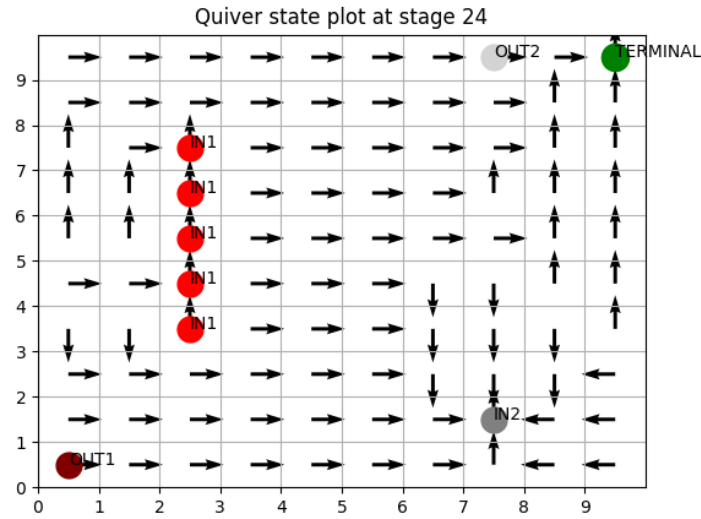
## 4 References

- Classroom lectures.
- Bertsekas: *Dynamic Programming and Optimal Control, Vol 2, 3rd ed.*
- Numpy, Matplotlib Dev Documentation.

Figure 1: Convergence of  $J_i$  till  $N = 10$  for **Terminal State (9,9)**Figure 2: Convergence of  $J_i$  till  $N = 25$  for **Terminal State (9,9)**

Figure 4: Convergence of  $J_i$  till  $N = 10$  for **Terminal State (3,0)**Figure 5: Convergence of  $J_i$  till  $N = 25$  for **Terminal State (3,0)**

Figure 7: Heat-map of  $J_i$  till  $N = 10$  for **Terminal State (9,9)**Figure 8: Quiver plot of  $\pi$  till  $N = 10$  for **Terminal State (9,9)**

Figure 9: Heat-map of  $J_i$  till  $N = 25$  for **Terminal State (9,9)**Figure 10: Quiver plot of  $\pi$  till  $N = 25$  for **Terminal State (9,9)**




Figure 11: Heat-map of  $J_i$  till absolute difference convergence for **Terminal State (9,9)**



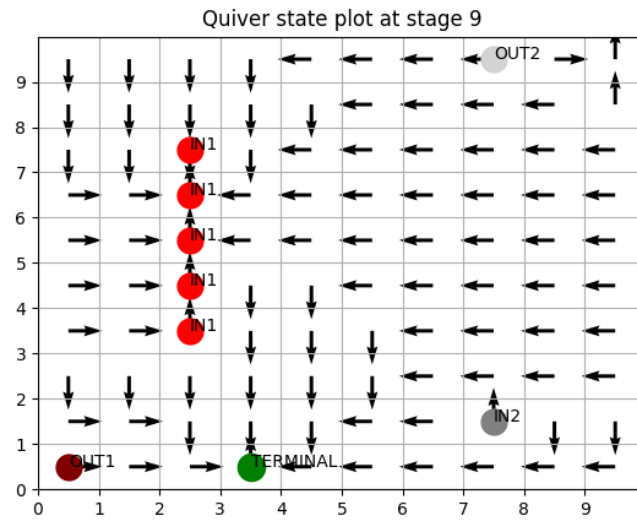


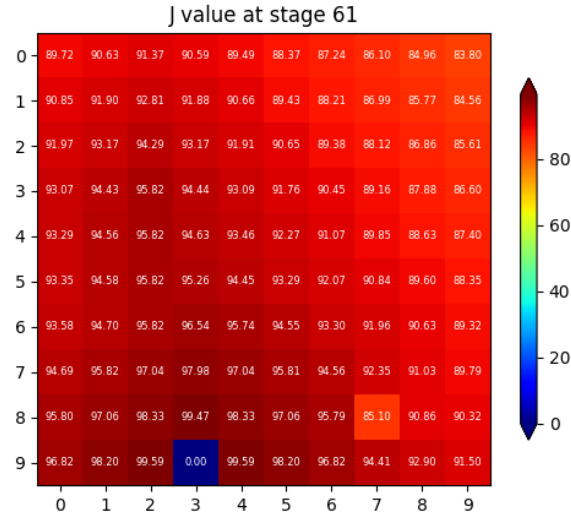
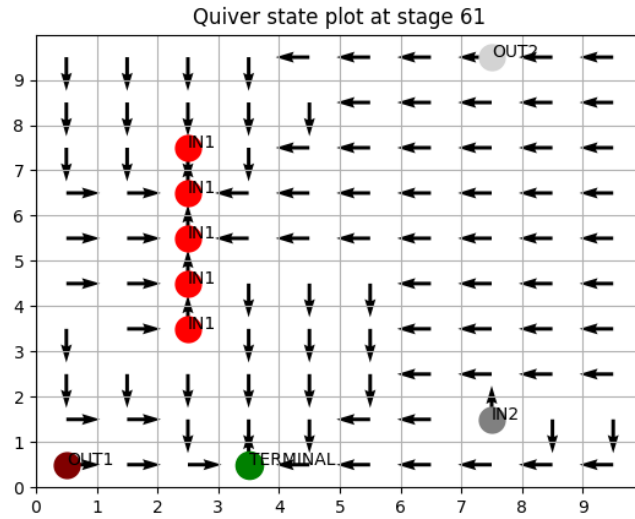


Figure 12: Quiver plot of  $\pi$  till absolute difference convergence for **Terminal State (9,9)**

Figure 13: Heat-map of  $J_i$  till  $N = 10$  for **Terminal State (3,0)**Figure 14: Quiver plot of  $\pi$  till  $N = 10$  for **Terminal State (3,0)**



Figure 15: Heat-map of  $J_i$  till  $N = 25$  for **Terminal State (3,0)**Figure 16: Quiver plot of  $\pi$  till  $N = 25$  for **Terminal State (3,0)**

Figure 17: Heat-map of  $J_i$  till absolute difference convergence for **Terminal State (3,0)**Figure 18: Quiver plot of  $\pi$  till absolute difference convergence for **Terminal State (3,0)**